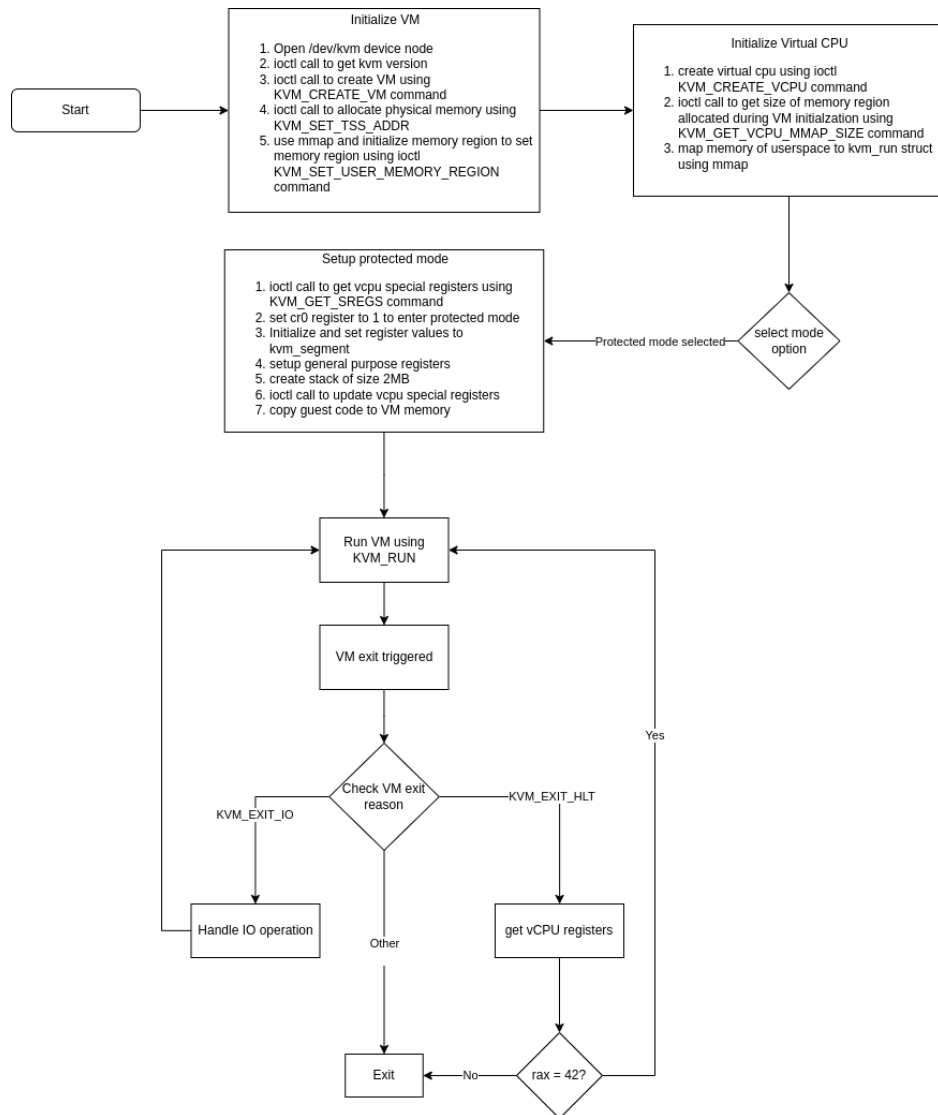


Part 1a report



1 Explanation of Code Snippets from `simple-kvm.c`

(a)

- `guest64[]` contains the program that will be executed by KVM when long mode is enabled.
- `guest64_end[]` marks the end of the program, and the program size is calculated as:

$$\text{guest64_end} - \text{guest64}$$

- So `guest64` and `guest64_end` act as start and end of a program.

(b)

- `PDE64_PRESENT` - Marks the page directory entry as present.
- `PDE64_RW` - marks entry as read/write enabled.
- `PDE64_USER` - marks page as user-mode accessible.
- `pdpt_addr, pd_addr` - marks physical address width limit.
- `PDE64_PS` - sets page size.
- `sregs->cr3 = pml4_addr;`
set `cr3` register with page map level 4 address which is a Page Directory Base Register.
- `sregs->cr4 = CR4_PAE;`
enables Physical address extension which expands physical address space to access memory beyond 4GB memory.
- `sregs->cr0 = CR0_PE | CR0_MP | CR0_ET | CR0_NE | CR0_WP | CR0_AM | CR0_PG;`
Enables:
 - Protected Mode
 - Monitor Co-Processor
 - Extension Type
 - Numeric Error
 - Write Protection
 - Alignment Mask

– Paging

- `sregs->efer = EFER_LME | EFER_LMA;`
sets bits in Extended Feature Enable Register(EFER) register to enable long mode(EFER_LME) and make long mode active(EFER_LMA).

(c)

```
vm->mem = mmap(NULL, mem_size, PROT_READ | PROT_WRITE,  
MAP_PRIVATE | MAP_ANONYMOUS | MAP_NORESERVE, -1, 0);
```

Allocates memory region of VM with read-write enabled.

MAP_PRIVATE - updates made to mapping are not visible to other process.

MAP_ANONYMOUS - mapping is not backed by any file.

MAP_NORESERVE - disable swap space reservation which does not allow mapping to be modified.

```
madvise(vm->mem, mem_size, MADV_MERGEABLE);
```

This tells kernel to enable Kernel Samepage Merging (KSM) for the pages in the range specified. it enables kernel to merge pages that have identical content into a single page.

(d)

- checks if the operation is IO_OUT by checking its direction and whether port used to write is 0xE9.
- gets the pointer to `kvm_run` which stores the data that is to be written.
- prints the data to stdout stream which is stored at an offset - `io.data_offset` from base address of `kvm_run` struct.

(e)

```
memcpy(&memval, &vm->mem[0x400], sz);
```

Copies value at memory location 0x400 of size sz into memval used in KVM_EXIT_HLT during exit to check if guest os outputs correct value.